

Multivariable logistic regression

GENERALIZED LINEAR MODELS IN PYTHON



Ita Cirovic Donev
Data Science Consultant

Multivariable setting

- Model formula

$$\text{logit}(y) = \beta_0 + \beta_1 x_1$$

Multivariable setting

- Model formula

$$\text{logit}(y) = \beta_0 + \beta_1 x_1$$

Multivariable setting

- Model formula

$$\text{logit}(y) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$$

Multivariable setting

- Model formula

$$\text{logit}(y) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$$

- In Python

```
model = glm('y ~ x1 + x2 + x3 + x4',  
            data = my_data,  
            family = sm.families.Binomial()).fit()
```

Example - well switching

```
formula = 'switch ~ distance100 + arsenic'
```

```
wells_fit = glm(formula = formula, data = wells,  
                family = sm.families.Binomial()).fit()
```

```
=====
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	0.0027	0.079	0.035	0.972	-0.153	0.158
distance100	-0.8966	0.104	-8.593	0.000	-1.101	-0.692
arsenic	0.4608	0.041	11.134	0.000	0.380	0.542

```
=====
```

Example - well switching

	coef	std err	z	P> z	[0.025	0.975]

Intercept	0.0027	0.079	0.035	0.972	-0.153	0.158
distance100	-0.8966	0.104	-8.593	0.000	-1.101	-0.692
arsenic	0.4608	0.041	11.134	0.000	0.380	0.542

- Both coefficients are statistically significant
- Sign of coefficients logical
- A unit-change in `distance100` corresponds to a negative difference of 0.89 in the logit
- A unit-change in `arsenic` corresponds to a positive difference of 0.46 in the logit

Impact of adding a variable

- Impact of `arsenic` variable
- `distance100` changes from -0.62 to -0.89
- Further away from the safe well
 - More likely to have higher arsenic levels

	coef	std err
Intercept	0.0027	0.079
distance100	-0.8966	0.104
arsenic	0.4608	0.041

	coef	std err
Intercept	0.6060	0.060
distance100	-0.6291	0.097

Multicollinearity

- Variables that are **correlated** with other model variables



- Increase in standard errors of coefficients
 - Coefficients may not be statistically significant

¹ https://en.wikipedia.org/wiki/Correlation_and_dependence

Presence of multicollinearity?

What to look for?

- Coefficient is not significant, but variable is highly correlated with y
- Adding/removing a variable significantly changes coefficients
- Not logical sign of the coefficient
- Variables have high pairwise correlation

Variance inflation factor (VIF)

- Most widely used diagnostic for multicollinearity
 - Computed for each explanatory variable
 - How inflated the variance of the coefficient is
- Suggested threshold $VIF > 2.5$
- In Python

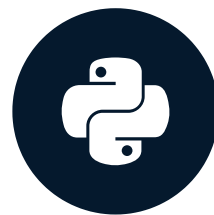
```
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

Let's practice!

GENERALIZED LINEAR MODELS IN PYTHON

Comparing models

GENERALIZED LINEAR MODELS IN PYTHON



Ita Cirovic Donev
Data Science Consultant

Deviance

- Formula

$$D = -2LL(\beta)$$

- Measure of error
- Lower deviance \rightarrow better model fit
- Benchmark for comparison is the **null deviance** \rightarrow intercept-only model
- Evaluate
 - Adding a random noise variable would, on average, decrease deviance by 1
 - Adding p predictors to the model deviance should decrease by more than p

Deviance in Python

Generalized Linear Model Regression Results						
=====						
Dep. Variable:	switch	No. Observations:	3020			
Model:	GLM	Df Residuals:	3018			
Model Family:	Binomial	Df Model:	1			
Link Function:	logit	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-2038.1			
Date:	Mon, 08 Apr 2019	Deviance:	4076.2			
Time:	10:24:56	Pearson chi2:	3.02e+03			
No. Iterations:	4	Covariance Type:	nonrobust			
=====						
	coef	std err	z	P> z	[0.025	0.975]

Intercept	0.6060	0.060	10.047	0.000	0.488	0.724
distance100	-0.6219	0.097	-6.383	0.000	-0.813	-0.431
=====						

Compute deviance

- Extract null-deviance and deviance

```
# Extract null deviance  
print(model.null_deviance)
```

```
4118.0992
```

```
# Extract model deviance  
print(model.deviance)
```

```
4076.2378
```

- Compute deviance using log likelihood

```
print(-2*model.llf)
```

```
4076.2378
```

- Reduction in deviance by 41.86
- Including `distance100` improved the fit

Model complexity

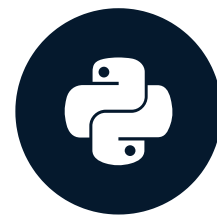
- `model_1` and `model_2` , where
 - $L1 > L2$
 - Number of parameters higher in `model_2`
- `model_2` is **overfitting**

Let's practice!

GENERALIZED LINEAR MODELS IN PYTHON

Model formula

GENERALIZED LINEAR MODELS IN PYTHON




Ita Cirovic Donev
Data Science Consultant

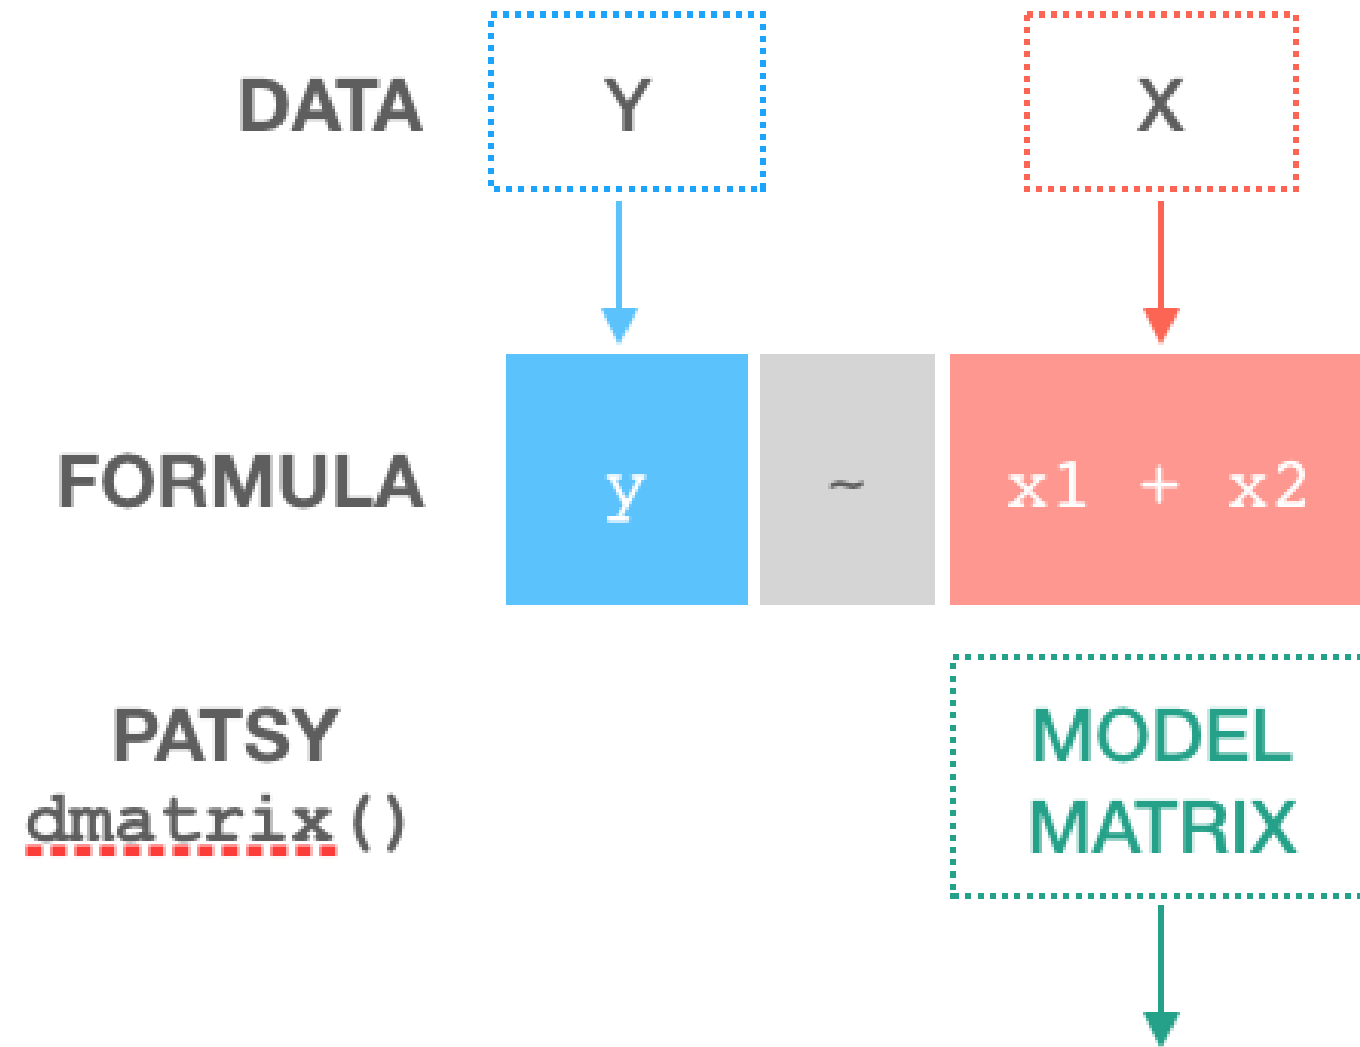
Formula and model matrix



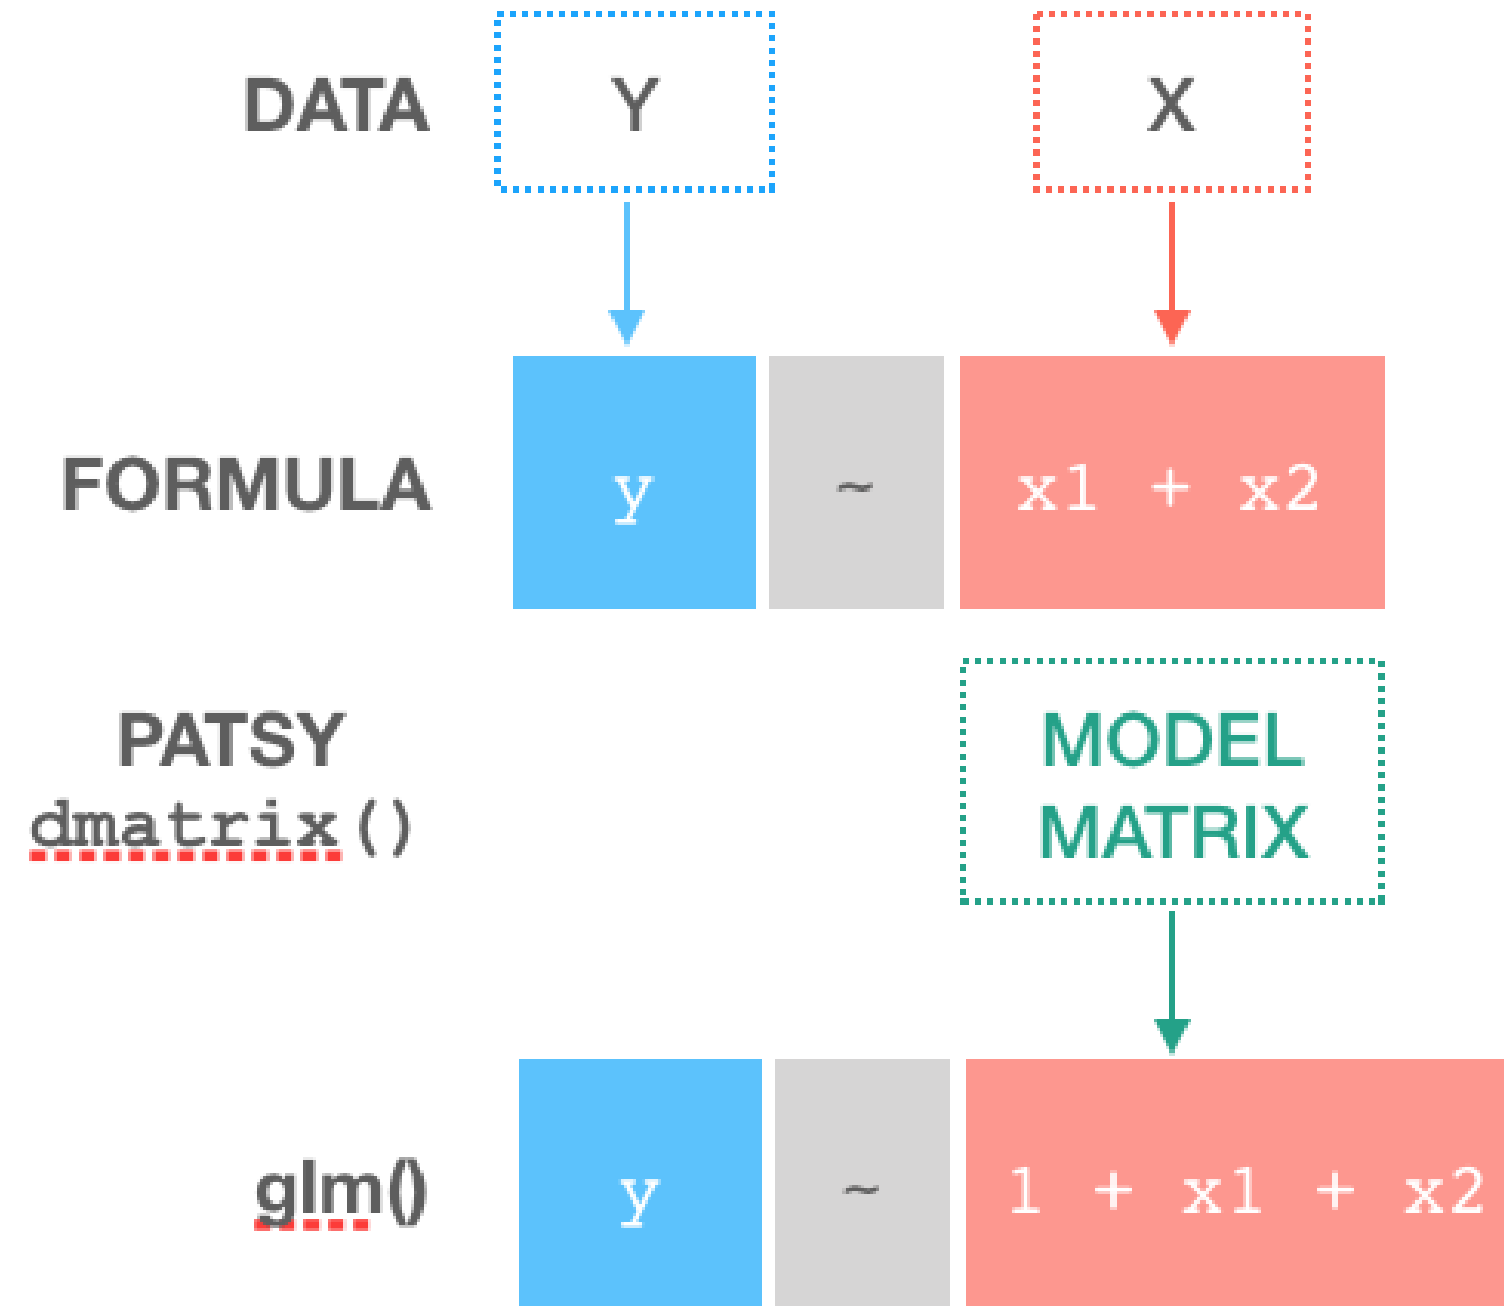
Formula and model matrix



Formula and model matrix



Formula and model matrix



Model matrix

- Model matrix: $y \sim \mathbf{X}$
- Model formula

```
'y ~ x1 + x2'
```

- Check model matrix structure

```
from patsy import dmatrix  
dmatrix('x1 + x2')
```

Intercept	x1	x2
1	1	4
1	2	5
1	3	6

Variable transformation

```
import numpy as np
```

```
'y ~ x1 + np.log(x2)'
```

```
dmatrix('x1 + np.log(x2)')
```

DesignMatrix with shape (3, 3)

Intercept	x1	np.log(x2)
1	1	1.38629
1	2	1.60944
1	3	1.79176

Centering and standardization

- Stateful transforms

```
'y ~ center(x1) + standardize(x2)'
```

```
dmatrix('center(x1) + standardize(x2)')
```

```
DesignMatrix with shape (3, 3)
  Intercept  center(x1)  standardize(x2)
        1         -1        -1.22474
        1          0         0.00000
        1          1         1.22474
```

Build your own transformation

```
def my_transformation(x):  
    return 4 * x
```

```
dmatrix('x1 + x2 + my_transformation(x2)')
```

DesignMatrix with shape (3, 4)

Intercept	x1	x2	my_transformation(x2)
1	1	4	16
1	2	5	20
1	3	6	24

Arithmetic operations

```
x1 = np.array([1, 2, 3])  
x2 = np.array([4, 5, 6])  
  
dmatrix('I(x1 + x2)')
```

```
DesignMatrix with shape (3, 2)  
  Intercept  I(x1 + x2)  
      1      5  
      1      7  
      1      9
```

```
x1 = [1, 2, 3]  
x2 = [4, 5, 6]  
  
dmatrix('I(x1 + x2)')
```

```
DesignMatrix with shape (6, 2)  
  Intercept  I(x1 + x2)  
      1      1  
      1      2  
      1      3  
      1      4  
      1      5  
      1      6
```

Coding the categorical data


Color
Red
Green
Blue

Coding the categorical data

Color	Color
Red	Green
Green	Red
Blue	Red
	Blue
	...

Coding the categorical data

Color	Color
Red	Green
Green	Red
Blue	Red
	Blue
	...



Red	Green	Blue
0	1	0
1	0	0
1	0	0
0	0	1
...

Patsy coding

- Strings and booleans are automatically coded
- Numerical \rightarrow categorical
 - `C()` function
- Reference group
 - Default: first group
 - `Treatment`
 - `levels`

The C() function

- Numeric variable

```
dmatrix('color', data = crab)
```

```
DesignMatrix with shape (173, 2)
```

```
Intercept  color
```

```
1         2
```

```
1         3
```

```
1         1
```

```
[... rows omitted]
```

- How many levels?

```
crab['color'].value_counts()
```

```
2    95
```

```
3    44
```

```
4    22
```

```
1    12
```

The C() function

- Categorical variable

```
dmatrix('C(color)', data = crab)
```

```
DesignMatrix with shape (173, 4)
```

Intercept	C(color)[T.2]	C(color)[T.3]	C(color)[T.4]
1	1	0	0
1	0	1	0
1	0	0	0

```
[... rows omitted]
```

Changing the reference group

```
dmatrix('C(color, Treatment(4))', data = crab)
```

```
DesignMatrix with shape (173, 4)
```

Intercept	C(color)[T.1]	C(color)[T.2]	C(color)[T.3]
1	0	1	0
1	0	0	1
1	1	0	0
[... rows omitted]			

Changing the reference group

```
l = [1, 2, 3, 4]
dmatrix('C(color, levels = l)', data = crab)
```

```
DesignMatrix with shape (173, 4)
  Intercept  C(color)[T.2]  C(color)[T.3]  C(color)[T.4]
           1             1             0             0
           1             0             1             0
           1             0             0             0
[... rows omitted]
```

Multiple intercepts

```
'y ~ C(color)-1'  
dmatrix('C(color)-1', data = crab)
```

```
DesignMatrix with shape (173, 4)  
  C(color)[1]  C(color)[2]  C(color)[3]  C(color)[4]  
           0           1           0           0  
           0           0           1           0  
           1           0           0           0  
[... rows omitted]
```

Let's practice!

GENERALIZED LINEAR MODELS IN PYTHON

Categorical and interaction terms

GENERALIZED LINEAR MODELS IN PYTHON



Ita Cirovic Donev
Data Science Consultant

Categorical variables

- Simple binary variable
 - Yes, No
- Nominal variables
 - Color: red, green, blue
- Ordinal variables
 - Levels of education: Education1, Education2,...,Education4

Analysis of covariance

- Explanatory variables
 - x_1 : categorical (binary)
 - x_2 : continuous
- Logistic model

$$\text{logit}(y = 1|X) = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

Analysis of covariance

- Explanatory variables
 - x_1 : categorical (binary)
 - x_2 : continuous
- Logistic model

$$\text{logit}(y = 1 | \textcolor{red}{X}) = \beta_0 + \beta_1 \textcolor{red}{x}_1 + \beta_2 \textcolor{red}{x}_2$$

Analysis of covariance

- Explanatory variables
 - x_1 : categorical (binary)
 - x_2 : continuous
- Logistic model

$$\text{logit}(y = 1|X) = \beta_0 + \beta_1 \mathbf{x_1} + \beta_2 x_2$$

- If $x_1 = 0$ then

$$\text{logit}(y = 1|\mathbf{x_1} = \mathbf{0}, x_2) = \beta_0 + \mathbf{0} + \beta_2 x_2$$

Analysis of covariance

- Explanatory variables
 - x_1 : categorical (binary)
 - x_2 : continuous
- Logistic model

$$\text{logit}(y = 1|X) = \beta_0 + \beta_1 \mathbf{x_1} + \beta_2 x_2$$

- If $x_1 = 0$ then

$$\text{logit}(y = 1|x_1 = 0, x_2) = \beta_0 + 0 + \beta_2 x_2$$


- If $x_1 = 1$ then

$$\text{logit}(y = 1|\mathbf{x_1 = 1}, x_2) = \beta_0 + \mathbf{\beta_1} + \beta_2 x_2$$

$$\text{logit}(y = 1|\mathbf{x_1 = 1}, x_2) = (\beta_0 + \mathbf{\beta_1}) + \beta_2 x_2$$


Assumptions

No interaction





Assumptions

No interaction



Assumptions

No interaction



Interactions

- Not equal slopes \rightarrow presence of **interaction**
- The effect of x_1 on y depends on the level of x_2 and vice versa
- Logistic model allowing for interactions

$$\text{logit}(y = 1|X) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2$$

Interactions

- Not equal slopes \rightarrow presence of **interaction**
- The effect of x_1 on y depends on the level of x_2 and vice versa
- Logistic model allowing for interactions

$$\text{logit}(y = 1|X) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2$$

- If $x_1 = 0$ then

$$\text{logit}(y = 1|x_1 = 0, x_2) = \beta_0 + 0 + \beta_2 x_2 + 0$$

Interactions

- Not equal slopes \rightarrow presence of **interaction**
- The effect of x_1 on y depends on the level of x_2 and vice versa

- Logistic model allowing for interactions

$$\text{logit}(y = 1|X) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2$$

- If $x_1 = 0$ then

$$\text{logit}(y = 1|x_1 = 0, x_2) = \beta_0 + \beta_2 x_2$$

- If $x_1 = 1$ then

$$\text{logit}(y = 1|x_1 = 1, x_2) = \beta_0 + \beta_1 + \beta_2 x_2 + \beta_3 x_2$$

$$\text{logit}(y = 1|x_1 = 1, x_2) = (\beta_0 + \beta_1) + (\beta_2 + \beta_3)x_2$$

Interactions

- Not equal slopes \rightarrow presence of **interaction**
- The effect of x_1 on y depends on the level of x_2 and vice versa

- Logistic model allowing for interactions

$$\text{logit}(y = 1|X) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2$$

- If $x_1 = 0$ then


$$\text{logit}(y = 1|x_1 = 0, x_2) = \beta_0 + \beta_2 x_2$$

- If $x_1 = 1$ then

$$\text{logit}(y = 1|x_1 = 1, x_2) = \beta_0 + \beta_1 + \beta_2 x_2 + \beta_3 x_2$$

$$\text{logit}(y = 1|x_1 = 1, x_2) = (\beta_0 + \beta_1) + (\beta_2 + \beta_3) x_2$$

Visualizing interactions



Interactions allow for:

- intercept and slope different for x_1
- β_1 : difference between the two intercepts
- β_3 : difference between the two slopes

Interaction types

- binary \times binary
- binary \times categorical
- binary \times continuous
- continuous \times categorical
- continuous \times continuous
- categorical \times categorical
- more than 2 variable interactions

Let's practice!

GENERALIZED LINEAR MODELS IN PYTHON

Congratulations!

GENERALIZED LINEAR MODELS IN PYTHON



Ita Cirovic Donev
Data Science Consultant

MODEL

- **Data** →
- **Link function** →
- **Model** →
- **1-unit increase in x** →

LINEAR MODEL

- Continuous
- Identity
- $y = \beta_0 + \beta_1 x_1$
- increases **y** by β_1

LOGISTIC REGRESSION

- Binary
- Logit
- $\text{logit}(y) = \beta_0 + \beta_1 x_1$
- increases **log odds** by β_1

POISSON REGRESSION

- Count
- Logarithm
- $\log(\lambda) = \beta_0 + \beta_1 x_1$
- **multiplies λ** by $\exp(\beta_1)$

MAIN PYTHON FUNCTIONS

- Fit the model
`statmodels` →

LINEAR MODEL

```
glm('y ~ x', data)
```

```
glm('y ~ x', data,  
    family = sm.families.Gaussian())
```

LOGISTIC REGRESSION

```
glm('y ~ x', data,  
    family = sm.families.Binomial())
```

POISSON REGRESSION

```
glm('y ~ x', data,  
    family = sm.families.Poisson())
```

Next steps...

- DataCamp courses
- Excellent reference books
 - *Regression Modeling Strategies* by Frank E. Harrell, Jr.
 - *An Introduction to Categorical Data Analysis* by Alan Agresti
 - *Applied Predictive Modeling* by Max Kuhn and Kjell Johnson

Happy modeling!

GENERALIZED LINEAR MODELS IN PYTHON