

## Step-2 Details Instructions:

### 1. List of variables selected for model built.

- Unemployment rate
- 10-year Treasury yield
- Prime rate',
- 'House Price Index (Level)'
- Unemployment rate\_lag\_1',
- 'Unemployment rate\_lag\_2
- Log of (Dow Jones Total Stock Market Index (Level)),
- House Price Index (Level)\_YOY
- Mortgage rate.

### 1. Correlation and Autocorrelation Requirements:

- For these nine variables create a correlation matrix and heatmap (table and graph). Multicollinearity for each variable, based on VIF values.
- Create a summary statistic, with number of observations, Mean, Std. Dev, Sum, Minimum and Maximum.
- Autocorrelation and White Noise Test: Perform Augmented Dickey-Fuller Unit Root Test (Rho, Tau, F values) to highlight autocorrelation. You can also visually inspect the autocorrelation function plot or perform statistical tests such as the Ljung-Box test or the Durbin-Watson test to check for significant autocorrelation.
- Exclude variables with negative results (you can send the results to me at this time if required to further narrow the list down).

### 1. For the remaining variables built models,

- I would recommend using the following (which I follow in SAS): GLIMMIX procedure in SAS, random effects (random residuals), binomial distribution, link equals logit.
- You can try other alternatives (as part of step C), but please add comments on relevant code so that I can understand the applied methods.
- model fit statistics should include AIC, BIC statistics and others ( I am assuming there is a command in Python which will generate all relevant one as there is in SAS). In-Sample Actual Default rate vs. Predicted default rate (curves).
- Perform seven to eight iterations using different combinations, such as Unemployment rate, 10-year Treasury yield.
- The results of each of these eight iterations, including the above-mentioned statistics and in-sample plots, should be added in the deliverable report.

1. Final variables (3 variables tops) should be based on variable signs, statistical significance, In-Sample RMSE, and other model fit statistics (BIC and AIC) of the eight iterations performed in step-2 above.

1. Once the model is finalized we will perform out-of-sample testing in the next steps.

**Note: Deliver a separate HTML, and a .py and an IPYNB files for this step. Please don't create this as an add-on to Step-1, instead a separate deliverable and files.**

```
In [1]: import pandas as pd
import numpy as np
```

```
In [2]: df_training = pd.read_excel('Datasets/Modeling Data-V03.xlsx', sheet_name='t')
df_training['DRS-Target Variable'] = df_training['DRS-Target Variable']/100
df_testing = pd.read_excel('Datasets/Modeling Data-V03.xlsx', sheet_name='t')
df_testing['DRS-Target Variable'] = df_testing['DRS-Target Variable']/100
```

```
In [3]: df_training
```

Out[3]:

	Scenario Name	Date	DRS-Target Variable	Real GDP growth	Nominal GDP growth	Real disposable income growth	Nominal disposable income growth	Unemployment rate	CPI inflation rate
0	Actual	2003 Q3	NaN	6.8	9.3	7.2	10.0	6.1	3.0
1	Actual	2003 Q4	NaN	4.7	7.3	1.1	3.1	5.8	1.5
2	Actual	2004 Q1	NaN	2.3	5.2	1.8	5.0	5.7	3.4
3	Actual	2004 Q2	NaN	3.2	6.5	4.2	7.1	5.6	3.2
4	Actual	2004 Q3	NaN	3.8	6.5	2.9	4.9	5.4	2.6
...	...	...	...	...	...	...	...	...	...
61	Actual	2018 Q4	0.0283	0.9	3.0	3.0	4.7	3.8	1.6
62	Actual	2019 Q1	0.0269	2.4	3.7	3.6	4.1	3.9	0.7
63	Actual	2019 Q2	0.0260	3.2	5.6	-1.4	1.3	3.6	3.5
64	Actual	2019 Q3	0.0244	2.8	4.1	2.3	3.4	3.6	1.3
65	Actual	2019 Q4	0.0234	1.9	3.6	2.4	4.1	3.6	2.6

66 rows × 19 columns

## Proprocessing

```
In [4]: def transformation(df,training=True,testing=True):
# Step 2: Perform data transformation - Log Transformation
# Defines a list called log_transform_variables
# that contains the names of variables to be log-transformed.

log_transform_variables = ['Dow Jones Total Stock Market Index (Level)',
                           'House Price Index (Level)',
                           'Commercial Real Estate Price Index (Level)']

# Loop through the variables to be log-transformed
# Applies the natural logarithm (np.log()) to the selected variable.
# Creates a new column with the log-transformed values using
for var in log_transform_variables:
    # Apply the natural logarithm to the selected variable and create a
    df[f'log_{var}'] = np.log(df[var])

# Step 3: Perform data transformation - Year-over-Year Change
# Defines a list called yoy_change_variables that contains the names
# of variables for which year-over-year changes will be calculated
yoy_change_variables = ['Dow Jones Total Stock Market Index (Level)',
                        'House Price Index (Level)',
                        'Commercial Real Estate Price Index (Level)']

# Loop through the variables for year-over-year change calculation
for var in yoy_change_variables:
    # Calculate the percentage change over a four-quarter period (assum
    # Creates a new column with the year-over-year change values using
    df[f'{var}_YOY'] = df[var].pct_change(3) * 100

# Step 4: Perform data transformation - Lags/Leads
# Defines the range of lags to be considered for lag/lead transformation
lags = range(1, 7) # Lags of up to six quarters
# Defines a list called lag_lead_variables that contains the names of v
lag_lead_variables = ['Unemployment rate', '10-year Treasury yield', 'M

# Loop through the variables for lag/lead transformation
# Loop through the variables for lag/lead transformation
for var in lag_lead_variables:
    # Loop through the specified lags
    for lag in lags:
        # Shift the variable values by the specified lag and create new
        # Shifts the variable values by the specified lag using
        df[f'{var}_lag_{lag}'] = df[var].shift(lag)

# this code will save the transformed data into csv file in the current
if training is True:
    df.to_csv('training-transformed_dataset.csv',index=False)
if testing is True:
    df = df.iloc[6:]
    df.to_csv('testing-transformed_dataset.csv',index=False)

return df

input_data = transformation(df_training,training=True)
testing_data = transformation(df_testing,testing=True)

print('shape of input Data :{}'.format(input_data.shape))
print('shape of test Data :{}'.format(testing_data.shape))

shape of input Data :(60, 43)
shape of test Data :(8, 43)
```

In [5]: `input_data.head()`

Out[5]:

	Scenario Name	Date	DRS-Target Variable	Real GDP growth	Nominal GDP growth	Real disposable income growth	Nominal disposable income growth	Unemployment rate	CPI inflation rate
6	Actual	2005 Q1	0.0142	4.5	7.9	-4.8	-2.5	5.3	2.0
7	Actual	2005 Q2	0.0155	2.0	5.0	3.9	6.6	5.1	2.7
8	Actual	2005 Q3	0.0159	3.2	7.0	1.7	6.1	5.0	6.2
9	Actual	2005 Q4	0.0164	2.3	5.6	3.4	6.7	5.0	3.8
10	Actual	2006 Q1	0.0160	5.5	8.5	8.3	10.6	4.7	2.1

5 rows × 43 columns

In [6]: `testing_data.head()`

Out[6]:

	Scenario Name	Date	DRS-Target Variable	Real GDP growth	Nominal GDP growth	Real disposable income growth	Nominal disposable income growth	Unemployment rate	CPI inflation rate
6	Actual	2020 Q1	0.0235	-5.1	-3.9	3.0	4.3	3.8	1.0
7	Actual	2020 Q2	0.0254	-31.2	-32.4	48.5	46.1	13.0	-3.1
8	Actual	2020 Q3	0.0284	33.8	38.7	-16.6	-13.6	8.8	4.7
9	Actual	2020 Q4	0.0274	4.5	6.6	-8.3	-6.9	6.8	2.4
10	Actual	2021 Q1	0.0267	6.3	10.9	54.7	60.6	6.2	3.7

5 rows × 43 columns

## Transformation

```
In [7]: def select_features(df):
        interested_features = [
            'DRS-Target Variable',
            '10-year Treasury yield',
            'Prime rate', 'House Price Index (Level)',
            'Unemployment rate_lag_1',
            'Unemployment rate_lag_2',
            'log_Dow Jones Total Stock Market Index (Level)',
            'House Price Index (Level)_YOY',
            'Mortgage rate',
        ]

        df = df[interested_features]
```

```
return df
```

```
In [8]: interested_features = [
        'DRS-Target Variable',
        '10-year Treasury yield',
        'Prime rate', 'House Price Index (Level)',
        'Unemployment rate_lag_1',
        'Unemployment rate_lag_2',
        'log_Dow Jones Total Stock Market Index (Level)',
        'House Price Index (Level)_YOY',
        'Mortgage rate',
        ]

training_data = select_features(input_data)
testing_data = select_features(testing_data)
```

```
In [ ]:
```

## Step-2

1. For these nine variables create a correlation matrix and heatmap (table and graph).  
Multicollinearity for each variable, based on VIF values.


```
In [9]: import seaborn as sns
import matplotlib.pyplot as plt

# This line calculates the correlation matrix of the df_transformed_trainin
# The correlation matrix is a square matrix that shows the correlation coef
# It provides a measure of the linear relationship between variables
correlation_matrix = training_data.corr()

# This line creates a heatmap using the Seaborn library.
# The heatmap() function is used to plot the correlation matrix as a color-
# The correlation_matrix is passed as the input data. The annot=True parame
# representing the correlation coefficients. The cmap='coolwarm' parameter
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')

# This line sets the title of the plot to 'Correlation Matrix Heatmap'
# using the title() function from the Matplotlib library
plt.title('Correlation Matrix Heatmap')
plt.show()

# Summary
# Given Code : calculates the correlation matrix for a given DataFrame and
# The heatmap provides a visual representation of the correlation between v
# with higher correlation values shown in warmer colors and lower correlati
```



```
In [10]: # This line imports the variance_inflation_factor function
# from the statsmodels.stats.outliers_influence module.
# This function is used to calculate the variance inflation factor,
# which is a measure of multicollinearity between variables in a regression
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Create a DataFrame to store the VIF values
# this line creates an empty DataFrame called vif_data to store
# the variable names and their corresponding VIF values
vif_data = pd.DataFrame()

# This line assigns the column names of the df_transformed_
# training DataFrame to the 'Variable' column in the vif_data
# DataFrame. This will store the names of the variables for
# which VIF values are calculated.
predictors_df = training_data.iloc[:,1:]
vif_data['Variable'] = predictors_df.columns

# This line calculates the VIF values for each variable
# in the df_transformed_training DataFrame and
# assigns them to the 'VIF' column in the vif_data DataFrame.
# The VIF values are computed using a list comprehension,
# where variance_inflation_factor is applied to each column of the
# df_transformed_training DataFrame using the range() function
vif_data['VIF'] = [variance_inflation_factor(predictors_df.values, i) for i

vif_data
```

```
# summary,
# Given code calculates the VIF values for each variable in the df_transformed
# The VIF values indicate the degree of multicollinearity between variables
# with higher values indicating stronger multicollinearity.
# The results are stored in a DataFrame called vif_data and
# then printed to the console.
```

Out[10]:

	Variable	VIF
0	10-year Treasury yield	399.922084
1	Prime rate	84.061378
2	House Price Index (Level)	1584.041360
3	Unemployment rate_lag_1	721.293720
4	Unemployment rate_lag_2	846.707158
5	log_Dow Jones Total Stock Market Index (Level)	2915.691100
6	House Price Index (Level)_YOY	3.165315
7	Mortgage rate	1066.915484

```
In [11]: updated_predictor = predictors_df.drop('House Price Index (Level)',axis=1)
updated_predictor
vif_data = pd.DataFrame()
vif_data['Variable'] = updated_predictor.columns
vif_data['VIF'] = [variance_inflation_factor(updated_predictor.values, i) for i in range(updated_predictor.shape[0])]
vif_data
```

Out[11]:

	Variable	VIF
0	10-year Treasury yield	283.130443
1	Prime rate	56.900745
2	Unemployment rate_lag_1	715.184761
3	Unemployment rate_lag_2	660.174255
4	log_Dow Jones Total Stock Market Index (Level)	141.840967
5	House Price Index (Level)_YOY	2.193561
6	Mortgage rate	730.413548

```
In [12]: import pandas as pd
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Drop the specified columns from the predictors_df
columns_to_drop = ['House Price Index (Level)', 'log_Dow Jones Total Stock Market Index (Level)']
updated_predictor = predictors_df.drop(columns_to_drop, axis=1)

# Calculate VIF for each variable in updated_predictor
vif_data = pd.DataFrame()
vif_data['Variable'] = updated_predictor.columns
vif_data['VIF'] = [variance_inflation_factor(updated_predictor.values, i) for i in range(updated_predictor.shape[0])]

# Print the VIF data
print(vif_data)
```

	Variable	VIF
0	10-year Treasury yield	76.471166
1	Prime rate	55.752033
2	Unemployment rate_lag_1	707.548535
3	Unemployment rate_lag_2	659.654433
4	House Price Index (Level)_YOY	1.621184
5	Mortgage rate	210.897864

```
In [13]: import pandas as pd
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Drop the specified columns from the predictors_df
columns_to_drop = ['Unemployment rate_lag_2', 'House Price Index (Level)',
updated_predictor = predictors_df.drop(columns_to_drop, axis=1)

# Calculate VIF for each variable in updated_predictor
vif_data = pd.DataFrame()
vif_data['Variable'] = updated_predictor.columns
vif_data['VIF'] = [variance_inflation_factor(updated_predictor.values, i) f

# Print the VIF data
print(vif_data)
```

	Variable	VIF
0	10-year Treasury yield	73.295917
1	Prime rate	50.581024
2	Unemployment rate_lag_1	18.334385
3	House Price Index (Level)_YOY	1.252499
4	Mortgage rate	209.631747

```
In [14]: import pandas as pd
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Drop the specified columns from the predictors_df
columns_to_drop = ['10-year Treasury yield', 'Unemployment rate_lag_2', 'Hou
updated_predictor = predictors_df.drop(columns_to_drop, axis=1)

# Calculate VIF for each variable in updated_predictor
vif_data = pd.DataFrame()
vif_data['Variable'] = updated_predictor.columns
vif_data['VIF'] = [variance_inflation_factor(updated_predictor.values, i) f

# Print the VIF data
print(vif_data)
```

	Variable	VIF
0	Prime rate	50.164525
1	Unemployment rate_lag_1	16.291920
2	House Price Index (Level)_YOY	1.236072
3	Mortgage rate	100.307276

```
In [15]: import pandas as pd
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Drop the specified columns from the predictors_df
columns_to_drop = ['Mortgage rate', '10-year Treasury yield', 'Unemployment r
updated_predictor = predictors_df.drop(columns_to_drop, axis=1)

# Calculate VIF for each variable in updated_predictor
vif_data = pd.DataFrame()
vif_data['Variable'] = updated_predictor.columns
vif_data['VIF'] = [variance_inflation_factor(updated_predictor.values, i) f

# Print the VIF data
print(vif_data)
```



	Variable	VIF
0	Prime rate	3.294323
1	Unemployment rate_lag_1	3.285024
2	House Price Index (Level)_Y0Y	1.051726

In [16]: `import statsmodels.api as sm`

```
# Perform Augmented Dickey-Fuller test for autocorrelation on each variable
# This line initializes an empty dictionary adf_results that will store the
# results of the Augmented Dickey-Fuller (ADF) test for each variable.
adf_results = {}

# This loop iterates over each column in the df DataFrame.
for column in training_data.columns:
    # This line performs the ADF test using the sm.tsa.stattools.adfuller()
    adf_test = sm.tsa.stattools.adfuller(training_data[column])
    # This line creates a pandas Series with the ADF test results and assign
    adf_results[column] = pd.Series(adf_test[:4], index=['ADF Statistic', '
# Display the ADF test results for each variable
print("\nAutocorrelation - Augmented Dickey-Fuller Test Results:")

# This loop iterates over each key-value pair in the adf_results dictionary
for column, results in adf_results.items():
    # This line prints the name of the variable being analyzed
    print("\nVariable:", column)
    # This line prints the ADF test results for the variable. The results v
    # p-value, number of lags used, and number of observations used
    print(results)

# Visualize autocorrelation function (ACF) plot for each variable

# This line calculates the number of variables in the DataFrame.
num_variables = len(training_data.columns)
# This line sets the number of columns for the subplots to be displayed
num_cols = 3 # Set the number of columns for subplots

# This line calculates the number of rows needed for the subplots based on
num_rows = (num_variables + num_cols - 1) // num_cols
# This line creates a figure and axes for the subplots, with the specified
# The figsize parameter determines the size of the figure.
fig, axes = plt.subplots(num_rows, num_cols, figsize=(15, num_rows*5))

# The following loop iterates over each variable in the DataFrame and plots
for i, column in enumerate(training_data.columns):
    # This line calculates the row index for the subplot based on the curre
    row = i // num_cols
    # : This line calculates the column index for the subplot based on the
    col = i % num_cols
    # This line selects the current subplot for plotting.
    ax = axes[row, col]

    # This line plots the autocorrelation function (ACF) for the variable
    sm.graphics.tsa.plot_acf(training_data[column], lags=len(training_data[
    # This line sets the title of the subplo
    ax.set_title("(ACF) Plot - " + column)
    # This line sets the label
    ax.set_xlabel("Lags")
    ax.set_ylabel("Autocorrelation")

# This line adjusts the layout of the subplots to prevent overlapping
plt.tight_layout()
```

```
# This line displays the subplots.  
plt.show()  
# In summary, this code performs the Augmented Dickey-Fuller (ADF) test for
```

## Autocorrelation - Augmented Dickey-Fuller Test Results:

Variable: DRS-Target Variable

ADF Statistic	-2.692612
p-value	0.075325
# Lags Used	4.000000
Number of Observations Used	55.000000
dtype:	float64

Variable: 10-year Treasury yield

ADF Statistic	-1.184460
p-value	0.680245
# Lags Used	2.000000
Number of Observations Used	57.000000
dtype:	float64

Variable: Prime rate

ADF Statistic	-3.469074
p-value	0.008818
# Lags Used	3.000000
Number of Observations Used	56.000000
dtype:	float64

Variable: House Price Index (Level)

ADF Statistic	-0.884130
p-value	0.793154
# Lags Used	5.000000
Number of Observations Used	54.000000
dtype:	float64

Variable: Unemployment rate\_lag\_1

ADF Statistic	-1.546206
p-value	0.510543
# Lags Used	1.000000
Number of Observations Used	58.000000
dtype:	float64

Variable: Unemployment rate\_lag\_2

ADF Statistic	-1.376505
p-value	0.593487
# Lags Used	1.000000
Number of Observations Used	58.000000
dtype:	float64

Variable: log\_Dow Jones Total Stock Market Index (Level)


ADF Statistic	-0.048824
p-value	0.954303
# Lags Used	0.000000
Number of Observations Used	59.000000
dtype:	float64

Variable: House Price Index (Level)\_YOY

ADF Statistic	-2.108353
p-value	0.241169
# Lags Used	2.000000
Number of Observations Used	57.000000
dtype:	float64

Variable: Mortgage rate

ADF Statistic	-1.092718
p-value	0.717967
# Lags Used	2.000000
Number of Observations Used	57.000000
dtype:	float64



```
In [17]: # Calculate summary statistics
summary_stats = training_data.describe().transpose()

# Display the summary statistics
print("\nSummary Statistics:")
summary_stats
```

Summary Statistics:

Out[17]:

	count	mean	std	min	25%	50%	75%
<b>DRS-Target Variable</b>	60.0	0.058122	0.033367	0.014200	0.028000	0.049450	0.093600
<b>10-year Treasury yield</b>	60.0	3.066667	1.040806	1.600000	2.200000	2.800000	3.900000
<b>Prime rate</b>	60.0	4.546667	1.674481	3.300000	3.300000	3.500000	5.325000
<b>House Price Index (Level)</b>	60.0	168.996667	24.676764	133.400000	142.775000	171.350000	190.225000
<b>Unemployment rate_lag_1</b>	60.0	6.131667	1.969642	3.600000	4.600000	5.300000	7.850000
<b>Unemployment rate_lag_2</b>	60.0	6.161667	1.943969	3.600000	4.675000	5.350000	7.850000
<b>log_Dow Jones Total Stock Market Index (Level)</b>	60.0	9.726728	0.353813	8.992707	9.469948	9.617690	9.982763
<b>House Price Index (Level)_YOY</b>	60.0	1.605091	5.982292	-13.422007	-1.425234	3.588308	4.408165
<b>Mortgage rate</b>	60.0	4.708333	0.989486	3.400000	3.900000	4.400000	5.725000

## Modeling the Data

In [18]: `columns_to_drop = ['Mortgage rate', 'Unemployment rate_lag_2', 'House Price Index (Level)_YOY']`  
`training_data.drop(columns_to_drop, axis=1, inplace=True)`  
`training_data.head(5)`

/tmp/ipykernel\_15112/2556610209.py:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`training_data.drop(columns_to_drop, axis=1, inplace=True)`

Out[18]:

	<b>DRS-Target Variable</b>	<b>10-year Treasury yield</b>	<b>Prime rate</b>	<b>Unemployment rate_lag_1</b>	<b>House Price Index (Level)_YOY</b>
<b>6</b>	0.0142	4.4	5.4	5.4	11.695906
<b>7</b>	0.0155	4.2	5.9	5.3	12.241055
<b>8</b>	0.0159	4.3	6.4	5.1	12.053301
<b>9</b>	0.0164	4.6	7.0	5.0	10.645724
<b>10</b>	0.0160	4.7	7.4	5.0	7.997763

In [19]: `columns_to_drop = ['Mortgage rate', 'Unemployment rate_lag_2', 'House Price Index (Level)_YOY']`  
`testing_data.drop(columns_to_drop, axis=1, inplace=True)`  
`testing_data.head(5)`

Out[19]:

	DRS-Target Variable	10-year Treasury yield	Prime rate	Unemployment rate_lag_1	House Price Index (Level)_YOY
6	0.0235	1.4	4.4	3.6	3.776291
7	0.0254	0.7	3.3	3.8	3.928064
8	0.0284	0.6	3.3	13.0	5.184493
9	0.0274	0.9	3.3	8.8	7.185629
10	0.0267	1.4	3.3	6.8	10.291439

Family Functions:

- `sm.families.Binomial()`: Binomial family for binary response data.
- `sm.families.Gaussian()`: Gaussian family for continuous response data.
- `sm.families.Poisson()`: Poisson family for count data.
- `sm.families.NegativeBinomial()`: Negative binomial family for count data.
- `sm.families.Gamma()`: Gamma family for continuous positive response data.
- `sm.families.InverseGaussian()`: Inverse Gaussian family for continuous positive response data.

Link Functions:

- `sm.families.links.logit()`: Logit link function for binary response data.
- `sm.families.links.identity()`: Identity link function for continuous response data.
- `sm.families.links.log()` or `sm.families.links.log1p()`: Log link function for count data (Poisson, Negative Binomial).
- `sm.families.links.inverse_power()`: Inverse power link function for Gamma and Inverse Gaussian families.

In addition to these options, statsmodels also provides support for custom family and link functions. ]

# Summary interpretations

To interpret the results of the GLM model, you can analyze various aspects of the model summary. Here are some key points to consider:

1. Coefficients: The coefficients indicate the estimated effect of each predictor on the response variable. They represent the average change in the response for a one-unit increase in the predictor, assuming all other predictors are held constant. For example, if the coefficient for "Prime rate" is 0.2, it suggests that, on average, a one-unit increase in the "Prime rate" is associated with a 0.2 increase in the "DRS-Target Variable."
2. Standard Errors: The standard errors provide an estimate of the variability or uncertainty associated with the coefficient estimates. Smaller standard errors indicate more precise estimates. In hypothesis testing, the standard errors are used to calculate t-statistics and p-values.

3. P-values: The p-values associated with the coefficients indicate the statistical significance of each predictor. They indicate the probability of observing a coefficient as extreme as the estimated coefficient if the null hypothesis (no effect) were true. Typically, a p-value below a certain threshold (e.g., 0.05) is considered statistically significant, suggesting a significant relationship between the predictor and the response.
4. AIC and BIC: The Akaike Information Criterion (AIC) and Bayesian Information Criterion (BIC) are measures of model fit that balance goodness of fit and model complexity. Lower AIC and BIC values indicate better-fitting models. You can compare the AIC and BIC values of different models to assess their relative quality.

Based on the provided code, you can examine the model summary output, including the coefficients, standard errors, p-values, AIC, and BIC. Interpretation of the coefficients depends on the specific dataset and context of your analysis. Remember to consider the scale and context of your variables when interpreting the coefficient values.

```
In [20]: import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt

# Load the dataset
# data = pd.DataFrame({
#     'DRS-Target Variable': [1.42, 1.55, 1.59, 1.64, 1.60],
#     'Prime rate': [5.4, 5.9, 6.4, 7.0, 7.4],
#     'Unemployment rate_lag_1': [5.4, 5.3, 5.1, 5.0, 5.0],
#     'House Price Index (Level)_YOY': [11.695906, 12.241055, 12.053301, 10
# })
# df_training = data
# Define the predictors and response variable
predictors = ['Prime rate', 'Unemployment rate_lag_1', 'House Price Index (L
response = 'DRS-Target Variable'

# Fit the GLM model
# 1 - sm.families.Gamma(sm.families.links.log())
# 2- sm.families.Gaussian(sm.families.links.identity())

model_1 = sm.GLM(training_data[response], sm.add_constant(training_data[pre
result_1 = model_1.fit()

# Fit the GLM with Gaussian family and identity link
model_2 = sm.GLM(training_data[response], sm.add_constant(training_data[pre
result_2 = model_2.fit()

# Fit the GLM with Tweedie family and appropriate link
model_3 = sm.GLM(training_data[response], sm.add_constant(training_data[pre
result_3 = model_3.fit()

# # Fit the GLM with Inverse Gaussian family and inverse squared link
# model_4 = sm.GLM(training_data[response], sm.add_constant(training_data[p
# result_4 = model_4.fit()

print('Summary of Model-1 : {}'.format(result_1.summary()))
print('Summary of Model-2 : {}'.format(result_2.summary()))
print('Summary of Model-3 : {}'.format(result_3.summary()))
```

```
# print('Summary of Model-4 : {}'.format(result_4.summary()))

print('AIC : {} and BIC : {} of Model-{}'.format(result_1.aic ,result_1.bic
print('AIC : {} and BIC : {} of Model-{}'.format(result_2.aic ,result_1.bic
print('AIC : {} and BIC : {} of Model-{}'.format(result_3.aic ,result_1.bic
# print('AIC : {} and BIC : {} of Model-{}'.format(result_4.aic ,result_1.b
```



Summary of Model-1 :  
Results

Generalized Linear Model Regression

```
=====
====
Dep. Variable:      DRS-Target Variable   No. Observations:
60
Model:              GLM                  Df Residuals:
56
Model Family:       Gamma                Df Model:
3
Link Function:      log                  Scale:              0.03
0485
Method:             IRLS                 Log-Likelihood:      20
1.51
Date:               Sat, 03 Jun 2023      Deviance:           1.
8416
Time:               11:01:27             Pearson chi2:
1.71
No. Iterations:     14                   Pseudo R-squ. (CS):
1.000
Covariance Type:    nonrobust
=====
```

		coef	std err	z	P> z
[0.025      0.975]					
-----					
-----					
const		-3.0124	0.155	-19.457	0.000
-3.316	-2.709				
Prime rate		-0.2163	0.017	-12.551	0.000
-0.250	-0.182				
Unemployment rate_lag_1		0.1646	0.015	11.180	0.000
0.136	0.193				
House Price Index (Level)_YOY		-0.0228	0.004	-5.809	0.000
-0.031	-0.015				

```
=====
====
Summary of Model-2 :  
Results
```

Generalized Linear Model Regression

```
=====
====
Dep. Variable:      DRS-Target Variable   No. Observations:
60
Model:              GLM                  Df Residuals:
56
Model Family:       Gaussian            Df Model:
3
Link Function:      identity            Scale:              5.5057
e-05
Method:             IRLS                 Log-Likelihood:      21
1.15
Date:               Sat, 03 Jun 2023      Deviance:           0.003
0832
Time:               11:01:27             Pearson chi2:        0.0
0308
No. Iterations:     3                   Pseudo R-squ. (CS):
1.000
Covariance Type:    nonrobust
=====
```

		coef	std err	z	P> z
[0.025      0.975]					
-----					
-----					

const		0.0136	0.007	2.061	0.039
0.001	0.026				
Prime rate		-0.0066	0.001	-9.055	0.000
-0.008	-0.005				
Unemployment rate_lag_1		0.0123	0.001	19.715	0.000
0.011	0.014				
House Price Index (Level)_YOY		-0.0006	0.000	-3.511	0.000
-0.001	-0.000				

Summary of Model-3 : Generalized Linear Model Regression Results

```

=====
====
Dep. Variable:      DRS-Target Variable      No. Observations:
60
Model:              GLM      Df Residuals:
56
Model Family:       Tweedie      Df Model:
3
Link Function:      Log      Scale:              0.006
3200
Method:              IRLS      Log-Likelihood:
nan
Date:                Sat, 03 Jun 2023      Deviance:              0.3
7398
Time:                11:01:27      Pearson chi2:
0.354
No. Iterations:      11      Pseudo R-squ. (CS):
nan
Covariance Type:     nonrobust
=====
=====

```

		coef	std err	z	P> z
[0.025	0.975]				
const		-2.9942	0.156	-19.172	0.000
-3.300	-2.688				
Prime rate		-0.2177	0.019	-11.332	0.000
-0.255	-0.180				
Unemployment rate_lag_1		0.1613	0.014	11.878	0.000
0.135	0.188				
House Price Index (Level)_YOY		-0.0163	0.004	-4.270	0.000
-0.024	-0.009				

```

=====
AIC : -395.0206531506095 and BIC : -227.44171787120763 of Model-1
AIC : -414.29562486331406 and BIC : -227.44171787120763 of Model-2
AIC : nan and BIC : -227.44171787120763 of Model-3

```

```

/home/iffi/anaconda3/lib/python3.9/site-packages/statsmodels/genmod/general
ized_linear_model.py:1799: FutureWarning: The bic value is computed using t
he deviance formula. After 0.13 this will change to the log-likelihood base
d formula. This change has no impact on the relative rank of models compare
d using BIC. You can directly access the log-likelihood version using the `
bic_llf` attribute. You can suppress this message by calling statsmodels.ge
nmod.generalized_linear_model.SET_USE_BIC_LLF with True to get the LLF-base
d version now or False to retain the deviance version.
warnings.warn(

```

## In-Sampling Evaluations

## Using Model-1

```
In [21]: from sklearn.metrics import mean_squared_error

# Define the predictors and response variable
predictors = ['Prime rate', 'Unemployment rate_lag_1', 'House Price Index (
response = 'DRS-Target Variable'

results_df = pd.DataFrame()

results_df['Predicted'] = result_1.predict(sm.add_constant(training_data[pr
results_df['Actual'] = training_data[response].tolist()

# Calculate MSE
mse = mean_squared_error(results_df['Predicted'], results_df['Actual'])


# Calculate RMSE
rmse = np.sqrt(mse)

print("MSE:", mse)
print("RMSE:", rmse)

plt.plot(results_df.index, results_df['Actual'], label='Actual',marker='o')
plt.plot(results_df.index, results_df['Predicted'], label='Predicted',marke
plt.xlabel('Observation')
plt.ylabel('DRS Target Variable')
plt.title('InSampling Results Curves')
plt.legend()
plt.show()
```

MSE: 0.00011939703862940322

RMSE: 0.010926895196230411



## Using Model-2

```
In [22]: from sklearn.metrics import mean_squared_error

# Define the predictors and response variable
predictors = ['Prime rate', 'Unemployment rate_lag_1', 'House Price Index (
response = 'DRS-Target Variable'

results_df = pd.DataFrame()

results_df['Predicted'] = result_2.predict(sm.add_constant(training_data[pr
results_df['Actual'] = training_data[response].tolist()

# Calculate MSE
mse = mean_squared_error(results_df['Predicted'], results_df['Actual'])

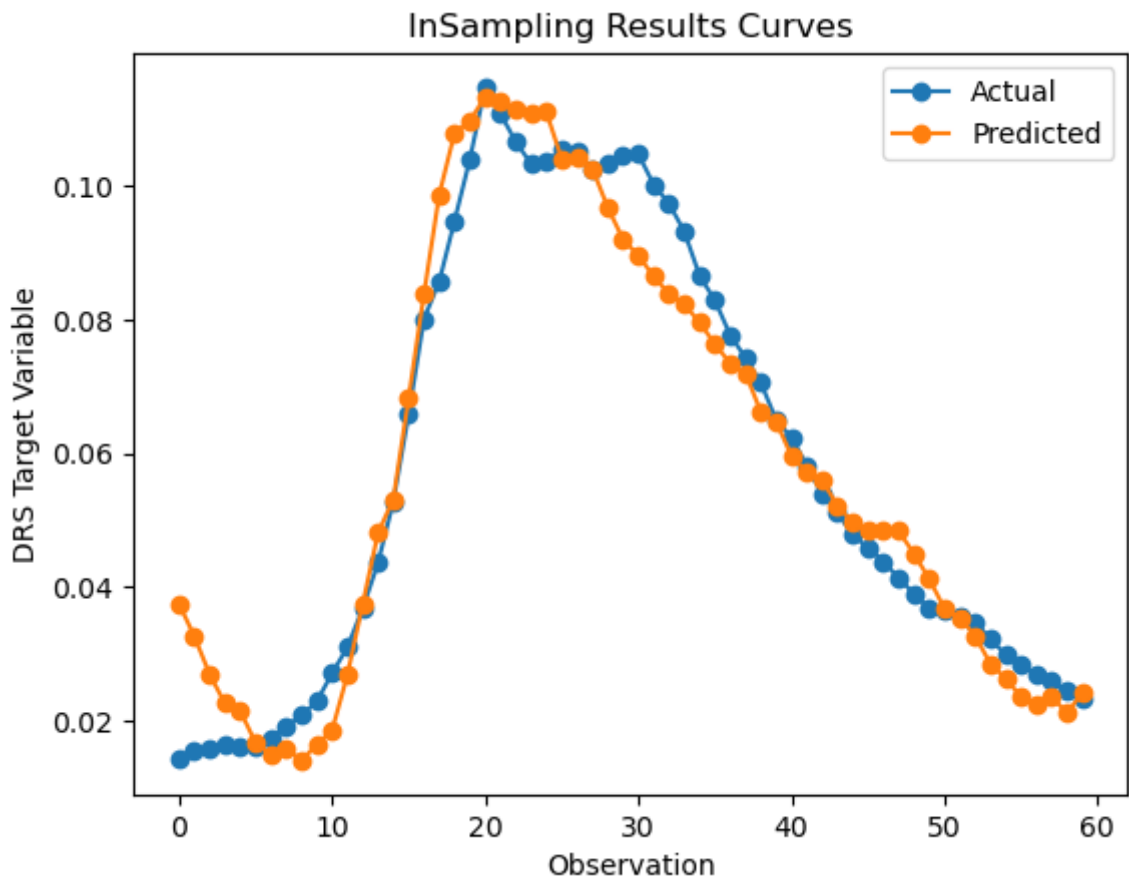
# Calculate RMSE
rmse = np.sqrt(mse)

print("MSE:", mse)
print("RMSE:", rmse)

plt.plot(results_df.index, results_df['Actual'], label='Actual',marker='o')
plt.plot(results_df.index, results_df['Predicted'], label='Predicted',marke
plt.xlabel('Observation')
plt.ylabel('DRS Target Variable')
plt.title('InSampling Results Curves')
plt.legend()
plt.show()
```

MSE: 5.138637590777052e-05

RMSE: 0.007168429110186591



## Using Model-3

```
In [23]: from sklearn.metrics import mean_squared_error

# Define the predictors and response variable
predictors = ['Prime rate', 'Unemployment rate_lag_1', 'House Price Index (
response = 'DRS-Target Variable'

results_df = pd.DataFrame()

results_df['Predicted'] = result_3.predict(sm.add_constant(training_data[pr
results_df['Actual'] = training_data[response].tolist()

# Calculate MSE
mse = mean_squared_error(results_df['Predicted'], results_df['Actual'])

# Calculate RMSE
rmse = np.sqrt(mse)

print("MSE:", mse)
print("RMSE:", rmse)

plt.plot(results_df.index, results_df['Actual'], label='Actual',marker='o')
plt.plot(results_df.index, results_df['Predicted'], label='Predicted',marke
plt.xlabel('Observation')
plt.ylabel('DRS Target Variable')
plt.title('InSampling Results Curves')
plt.legend()
plt.show()
```

MSE: 8.749385011173087e-05

RMSE: 0.009353814735803295

